

CSC485A/586A: Software Defined Networking

Assignment 1

Submission Date: May 16, 2017 on connex

As mentioned in the class, mininet (<http://mininet.org/>) is an emulation software that creates a **realistic virtual network**, running **real kernel, switch and application code**, on a single machine (VM, cloud or native). We will use this tool for a few exercises in this course. Hence, the purpose of this assignment is to get you going, use it to write a simple network topology and test it. As this is a 4th year/grad course, it is expected that you should be able to download this tool and figure it out on how to use it. There are lots of examples and help on the internet.

1. Please visit the mininet webpage (<http://mininet.org/>) where you can find instructions on how to get started (<http://mininet.org/download/>)
2. It is best to use Virtual Box (<https://www.virtualbox.org/>), a Type 2 hypervisor as it is free and works well on Windows, Mac and Linux machines.
3. Then download a mininet VM image from <https://github.com/mininet/mininet/wiki/Mininet-VM-Images>
4. Follow the Mininet walkthrough (<http://mininet.org/walkthrough/>) on how to use mininet, and familiarize with some of the commands and examples.
5. You will use “iperf” tool in order to study link bandwidth, delay and packet loss in networks. The mininet VM already comes with “iperf” tool.

Please note that each student must do these assignments individually. Students will get a “0” if the code is plagiarized or copied from other students. You can submit your code on connex. I will select students randomly to see their actual work. If students are not able to demonstrate their work when asked, they will be automatically assigned a “0”.

The following is an exercise to build custom topologies. This material is adopted from on-line resources by Prof. Nick Feamster (<http://noise.gatech.edu/classes/cs8803sdn/fall2014/>). Read this, understand and test the code in your mininet setup. The actual assignment is on Page 7. Please submit your solution back on connex.

In this exercise, you will be learning how to build custom topologies using Mininet Python API and how certain parameters like bandwidth, delay, loss and queue size can be set individually for different links in the topology. You’ll also learn how to do performance testing of these custom topologies using ping and iperf. After the overview, you will be asked to create and submit your own custom topology based on the most common 3-tier Datacenter architecture i.e., core, aggregation and edge. Please follow each step carefully.

Overview

The network you'll use in this exercise includes hosts and switches connected in a linear topology, as shown in the figure below.




Figure 1: Hosts and Switches connected in a linear topology

Creating Topology

Mininet supports *parametrized topologies*. With a few lines of Python code, you can create a flexible topology, which can be configured based on the parameters you pass to it, and reused for multiple experiments. For example, here is a simple network topology (based on Figure 1) which consists of a specified number of hosts (h1 through hN) connected to their individual switches (s1 through sN):

Linear Topology (without Performance Settings) (File name on connex: LinearTopo.py)

```
#!/usr/bin/python

from mininet.topo import Topo
from mininet.net import Mininet
from mininet.util import irange, dumpNodeConnections
from mininet.log import setLogLevel

class LinearTopo(Topo):
    "Linear topology of k switches, with one host per switch."

    def __init__(self, k=2, **opts):
```

```

    """Init.
        k: number of switches (and hosts)
        hconf: host configuration options
        lconf: link configuration options"""

    super(LinearTopo, self).__init__(**opts)

    self.k = k

    lastSwitch = None
    for i in xrange(1, k):
        host = self.addHost('h%s' % i)
        switch = self.addSwitch('s%s' % i)
        self.addLink( host, switch)
        if lastSwitch:
            self.addLink( switch, lastSwitch)
        lastSwitch = switch

def simpleTest():
    "Create and test a simple network"
    topo = LinearTopo(k=4)
    net = Mininet(topo)
    net.start()
    print "Dumping host connections"
    dumpNodeConnections(net.hosts)
    print "Testing network connectivity"
    net.pingAll()
    net.stop()

if __name__ == '__main__':
    # Tell mininet to print useful information
    setLogLevel('info')

    simpleTest()

```

Figure 1. LinearTopo.py

The important classes, methods, functions and variables in the above code include:

- **Topo**: the base class for Mininet topologies
- **addSwitch()**: adds a switch to a topology and returns the switch name
- **addHost()**: adds a host to a topology and returns the host name
- **addLink()**: adds a bidirectional link to a topology (and returns a link key, but this is not important). Links in Mininet are bidirectional unless noted otherwise.
- **Mininet**: main class to create and manage a network
- **start()**: starts your network
- **pingAll()**: tests connectivity by trying to have all nodes ping each other
- **stop()**: stops your network
- **net.hosts**: all the hosts in a network
- **dumpNodeConnections()**: dumps connections to/from a set of nodes.
- **setLogLevel('info' | 'debug' | 'output')**: set Mininet's default output level; 'info' is recommended as it provides useful information.

Additional example code may be found in [mininet/examples](https://github.com/mininet/mininet/tree/master/examples).

Setting Performance Parameters

In addition to basic behavioral networking, Mininet provides performance limiting and isolation features, through the `CPULimitedHost` and `TCLink` classes. There are multiple ways that these classes may be used, but one simple way is to specify them as the default host and link classes/constructors to `Mininet()`, and then to specify the appropriate parameters in the topology.

Linear Topology (with Performance Settings) (file name on conneX: `LinearTopoperf.py`)

```
#!/usr/bin/python
```

```
from mininet.topo import Topo
from mininet.net import Mininet
from mininet.node import CPULimitedHost
from mininet.link import TCLink
from mininet.util import irange, dumpNodeConnections
from mininet.log import setLogLevel
```

```
class LinearTopo(Topo):
```

```
    "Linear topology of k switches, with one host per switch."
```

```
    def __init__(self, k=2, **opts):
```

```
        """Init.
```

```
        k: number of switches (and hosts)
```

```
        hconf: host configuration options
```

```
        lconf: link configuration options"""
```

```
        super(LinearTopo, self).__init__(**opts)
```

```
        self.k = k
```

```
        lastSwitch = None
```

```
        for i in irange(1, k):
```

```
            host = self.addHost('h%s' % i, cpu=.5/k)
```

```
            switch = self.addSwitch('s%s' % i)
```

```
            # 10 Mbps, 5ms delay, 1% loss, 1000 packet queue
```

```
            self.addLink(host, switch, bw=10, delay='5ms', loss=1,
max_queue_size=1000, use_htb=True)
```

```
            if lastSwitch:
```

```
                self.addLink(switch, lastSwitch, bw=10, delay='5ms', loss=1,
max_queue_size=1000, use_htb=True)
```

```
            lastSwitch = switch
```

```
    def perfTest():
```

```
        "Create network and run simple performance test"
```

```

topo = LinearTopo(k=4)
net = Mininet(topo=topo,
               host=CPULimitedHost, link=TCLink)
net.start()
print "Dumping host connections"
dumpNodeConnections(net.hosts)
print "Testing network connectivity"
net.pingAll()
print "Testing bandwidth between h1 and h4"
h1, h4 = net.get('h1', 'h4')
net.iperf((h1, h4))
net.stop()

if __name__ == '__main__':
    setLogLevel('info')

    perfTest()

```

Some important methods and parameters:

`self.addHost(name, cpu=f)`: This allows you to specify a fraction of overall system CPU resources which will be allocated to the virtual host.

`self.addLink(node1, node2, bw=10, delay='5ms', max_queue_size=1000, loss=1, use_htb=True)`: adds a bidirectional link with bandwidth, delay and loss characteristics, with a maximum queue size of 1000 packets using the Hierarchical Token Bucket rate limiter and netem delay/loss emulator. The parameter `bw` is expressed as a number in Mb/s; `delay` is expressed as a string with units in place (e.g. '5ms', '100us', '1s'); `loss` is expressed as a percentage (between 0 and 100); and `max_queue_size` is expressed in packets.

You may find it useful to create a Python dictionary to make it easy to pass the same parameters into multiple method calls, for example:

```

linkopts = dict(bw=10, delay='5ms', loss=1, max_queue_size=1000,
                use_htb=True)

linkopts = dict(bw=10, delay='5ms', loss=1, max_queue_size=1000,
                use_htb=True)
'''
alternately: linkopts = {'bw':10, 'delay':'5ms', 'loss':1,
                        'max_queue_size':1000, 'use_htb':True}
'''
self.addLink(node1, node2, **linkopts)

```

Running in Mininet

To run the custom topology you have created above, follow the instructions below:

- Create a `LinearTopo.py` script on your Mininet VM and copy the contents of Linear Topology (without Performance Settings), listed above in it.

- Make the script executable
 - `$ chmod u+x LinearTopo.py`
- Execute the script
 - `$ sudo ./LinearTopo.py`
 -

Output

```
*** Creating network
```

```
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2 s3 s4
*** Adding links:
(h1, s1) (h2, s2) (h3, s3) (h4, s4) (s1, s2) (s2, s3) (s3, s4)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
*** Starting 4 switches
s1 s2 s3 s4
Dumping host connections
h1 h1-eth0:s1-eth1
h2 h2-eth0:s2-eth1
h3 h3-eth0:s3-eth1
h4 h4-eth0:s4-eth1
Testing network connectivity
*** Ping: testing ping reachability
h1 -> h2 h3 h4
h2 -> h1 h3 h4
h3 -> h1 h2 h4
h4 -> h1 h2 h3
*** Results: 0% dropped (0/12 lost)
*** Stopping 4 hosts
h1 h2 h3 h4
*** Stopping 4 switches
s1 ...s2 ....s3 .....s4 ...
*** Stopping 1 controllers
c0
*** Done
```

Assignment

Background

Data center networks typically have a tree-like topology. End-hosts connect to top-of-rack switches, which form the leaves (edges) of the tree; one or more core switches form the root; and one or more layers of aggregation switches form the middle of the tree. In a basic tree topology, each switch (except the core switch) has a single parent switch. Additional switches and links may be added to construct more complex tree topologies (e.g., fat tree) in an effort to improve fault tolerance or increase inter-rack bandwidth.

In this assignment, your task is to create a simple tree topology. You will assume each level i.e., core, aggregation, edge and host to be composed of a single layer of switches/hosts with a configurable fanout value (k). For example, a simple tree network having a single layer per each level and a fanout of 2 looks like:




Figure 2: Simple Tree Topology with Fanout 2

To start this exercise, download CustomTopo.py

- CustomTopo.py: a skeleton class, which you will update with the logic for creating the datacenter topology, described above.

CustomTopo.py

The skeleton class takes following arguments as input:

- `linkopts1`: for specifying performance parameters for the links between core and aggregation switches.
- `linkopts2`: for specifying performance parameters for the links between aggregation and edge switches.
- `linkopts3`: for specifying performance parameters for the links between edge switches and host
- `Fanout`: to specify fanout value i.e., number of childs per node.

Your logic should support setting at least bw and delay parameters for each link.

* These instructions are adapted from mininet.org and [wisc-cs838](#) and other internet resources.